

High performant in-memory datasets with Netflix Hollow

Who am I?

Roberto Perez Alcolea

- Mexican
- Streaming Platform @ Target
- Groovy Enthusiast
- roberto@perezalcolea.info
- @rpalcolea

Dataset Distribution

The problem

- Dissemination of small or moderately sized data sets (no big-data)

Common approaches:

- Sending data to a data store (RDMS, NoSQL)
- Serializing and keeping a local copy

Is there a solution?

NETFLIX | OSS H0110W

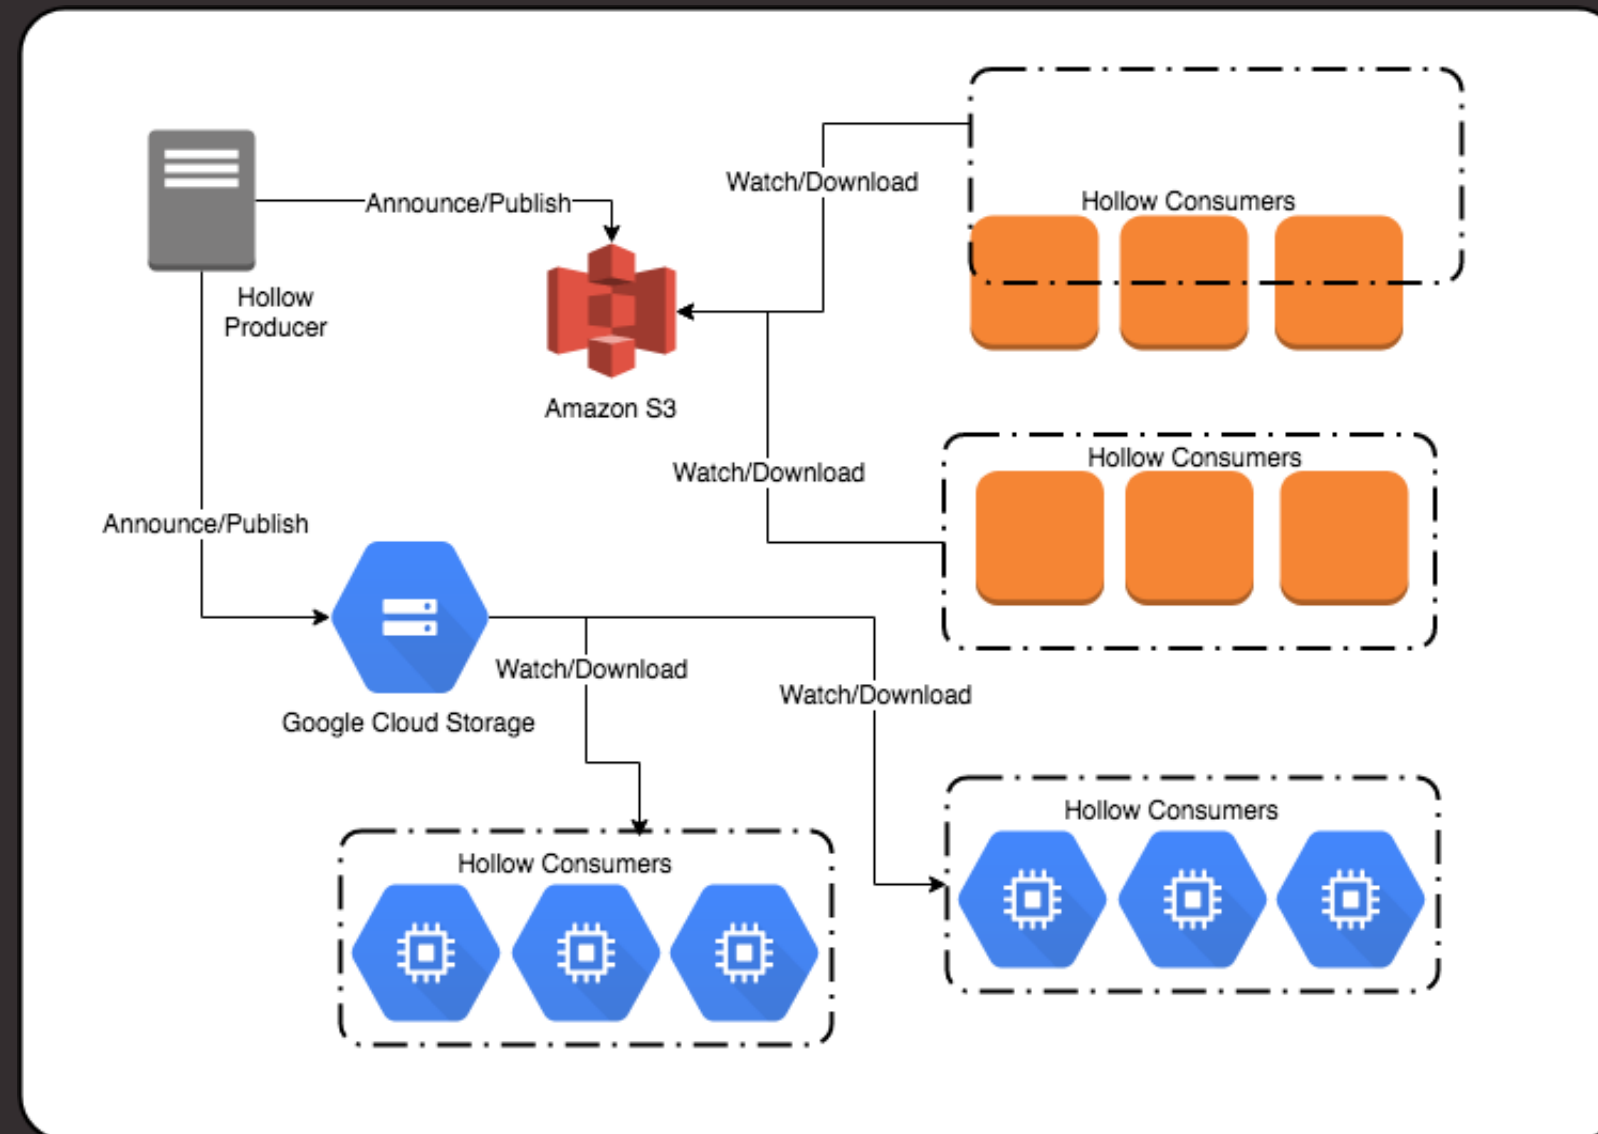
What is Hollow?

- Java library and toolset for disseminating in-memory datasets from a single producer to many consumers

Goals:

- Maximum development agility
- Highly optimized performance and resource management
- Extreme stability and reliability

How it works



Key concepts

- Blob / Blob Store
- Snapshot/Delta
- Type
- Deduplication

Key concepts

- State Engine
- Ordinal
 - unique identifier of the record within a type
 - sufficient to locate the record within a type
 - Immutability
- Cycle

Hollow Producer

- A single machine that retrieves all data from a source of truth and produces a delta chain
- Encapsulates the details of compacting, publishing, announcing, validating, and (if necessary) rollback of data states

Hollow Producer

```
class MyProducer {
    HollowProducer buildProducer() {
        ...
        HollowProducer
            .withPublisher(publisher)           /// required: a BlobPublisher
            .withAnnouncer(announcer)           /// optional: an Announcer
            .withValidators(validators)         /// optional: one or more Validator
            .withListeners(listeners)           /// optional: one or more HollowProducerListeners
            .withBlobStagingDir(dir)             /// optional: a java.io.File
            .withBlobCompressor(compressor)     /// optional: a BlobCompressor
            .withBlobStager(stager)             /// optional: a BlobStager
            .withSnapshotPublishExecutor(e)     /// optional: a java.util.concurrent.Executor
            .withNumStatesBetweenSnapshots(n)   /// optional: an int
            .withTargetMaxTypeShardSize(size)   /// optional: a long
            .build()
    }
}
```

Hollow Producer - Capabilities

- Restore at Startup
- Rolling Back
- Validating Data
- Compacting Data

Hollow Consumer

- Encapsulates the details of initializing and keeping a dataset up to date
 - At initialization time, loads snapshot
 - After initialization time, keeps a local copy of the dataset current by applying delta transitions
- Each time a new version is announced, `triggerRefresh()` should be called on the `HollowConsumer`

Hollow Consumer

```
class MyConsumer {
    HollowConsumer buildConsumer() {
        ...
        HollowConsumer
            .withBlobRetriever(blobRetriever)           /// required: a BlobRetriever
            .withLocalBlobStore(localDiskDir)           /// optional: a local disk location
            .withAnnouncementWatcher(announcementWatcher) /// optional: a AnnouncementWatcher
            .withRefreshListener(refreshListener)        /// optional: a RefreshListener
            .withGeneratedAPIClass(MyGeneratedAPI.class) /// optional: a generated client API class
            .withFilterConfig(filterConfig)              /// optional: a HollowFilterConfig
            .withDoubleSnapshotConfig(doubleSnapshotCfg) /// optional: a DoubleSnapshotConfig
            .withObjectLongevityConfig(objectLongevityCfg) /// optional: an ObjectLongevityConfig
            .withObjectLongevityDetector(detector)       /// optional: an ObjectLongevityDetector
            .withRefreshExecutor(refreshExecutor)        /// optional: an Executor
            .build()
    }
}
```


Hollow Consumer API Generation

- We can initialize the data model using our POJOs

```
HollowAPIGenerator generator = new HollowAPIGenerator.Builder().withAPIClassname("BooksAPI")
    .withPackageName("io.perezalcolea.hollow.api")
    .withDataModel(Book.class)
    .withDestination("./data-model/src/main/java")
    .build();

generator.generateSourceFiles();
```

Insight Tools

Hollow Explorer

- UI which can be used to browse and search records within any dataset

```
class MyExplorer {  
    void startExplorer() {  
        //Initialize consumer  
        HollowConsumer hollowConsumer = HollowConsumerBuilder.build("publish-dir", BooksAPI.class)  
        hollowConsumer.triggerRefresh()  
        HollowExplorerUIServer server = new HollowExplorerUIServer(hollowConsumer, 8080)  
        server.start()  
        server.join()  
    }  
}
```

Hollow History

- UI which can be used to browse and search changes in a dataset over time

```
class MyHistory {  
    void startHistory() {  
        //Initialize consumer  
        HollowConsumer hollowConsumer = HollowConsumerBuilder.build("publish-dir", BooksAPI.class)  
        hollowConsumer.triggerRefresh()  
        HollowHistoryUIServer server = new HollowHistoryUIServer(hollowConsumer, 8090)  
        server.start()  
        server.join()  
    }  
}
```

Heap Usage Analysis

- Given a loaded HollowReadStateEngine, it is possible to iterate over each type and gather statistics about its approximate heap usage

```
class MyMemoryInsight {  
    void printMemoryUsage() {  
        HollowReadStateEngine stateEngine = consumer.getStateEngine()  
  
        long totalApproximateHeapFootprint = 0  
  
        for(HollowTypeReadState typeState : stateEngine.typeStates) {  
            String typeName = typeState.schema.name  
            long heapCost = typeState.approximateHeapFootprintInBytes  
            println(typeName + ": " + heapCost);  
            totalApproximateHeapFootprint += heapCost  
        }  
  
        println("TOTAL: " + totalApproximateHeapFootprint)  
    }  
}
```

Other tools

- Metrics Collector
- Blob Storage Cleaner
- Filtering
- Combining/Splitting
- Patching

Interacting with a Hollow Dataset

Sample model

```
@HollowPrimaryKey(fields={"bookId"})
public class Book {
    long bookId;
    String title;
    String isbn;
    String publisher;
    String language;
    Set<Author> authors;
}

public class Author {
    long id;
    String authorName;
}
```


Indexing/Querying

Default Primary Keys

- Each type in our data model gets a custom index class called `<typename>PrimaryIndex`
- Backed by Hollow Consumer
- Will automatically stay up-to-date as your dataset updates

```
class MyPrimaryIndex {  
    Book findBook(long bookId) {  
        HollowConsumer hollowConsumer = //my consumer  
        BookPrimaryIndex bookPrimaryIndex = new BookPrimaryIndex(consumer)  
        Book book = bookPrimaryIndex.findMatch(bookId)  
    }  
}
```

Consumer-specified Primary Keys

- A primary key index is not restricted to just default primary keys

```
class MyPrimaryKeyIndex {  
    Book findAuthor(long authorId) {  
        HollowConsumer hollowConsumer = //my consumer  
        AuthorPrimaryIndex authorPrimaryIndex = new AuthorPrimaryIndex(consumer, "id")  
        Author author = authorPrimaryIndex.findMatch(authorId)  
    }  
}
```

Hash Index

- Records based on keys for which there is not a one-to-one mapping between records and key values
- Must specify each of a query type, a select field, and one or more match fields

```
class MyHashIndex {  
    Iterable<Book> findBooksByPublisher(String publisher) {  
        HollowConsumer hollowConsumer = //my consumer  
        BooksAPIHashIndex publisherHashIndex = new BooksAPIHashIndex(consumer, "Book", "", "publisher.value")  
        return publisherHashIndex.findMatch(publisher)  
    }  
}
```

Data Modeling

Primary Keys

- @HollowPrimaryKey annotation
- Provide a shortcut when creating a primary key index

```
@HollowPrimaryKey(fields={"bookId"})
public class Book {
    long bookId;
    String title;
    String isbn;
    String publisher;
    String language;
    Set<Author> authors;
}
```

Inline vs Referenced Fields

- Inline: fields that they are no longer REFERENCE fields, but instead encode their data directly in each record

```
@HollowPrimaryKey(fields={"bookId"})  
public class Book {  
    long bookId;  
    @HollowInline  
    String title;  
    String isbn;  
    String publisher;  
    String language;  
    Set<Author> authors;  
}
```

Namespaced Record Type Names

- Fields with like values may reference the same record type, but reference fields of the same primitive type elsewhere in the data model use different record types
- Types can be filtered

```
@HollowPrimaryKey(fields={"bookId"})
public class Book {
    long bookId;
    ...
    @HollowTypeName(name="Publisher")
    String publisher;
    @HollowTypeName(name="Language")
    String language;
    ...
}
```


Maintaining Backwards Compatibility

- Adding a new type
- Removing an existing type
- Adding a new field to an existing type
- Removing an existing field from an existing type.

Demo

Useful links

<https://hollow.how/>

<https://github.com/Netflix/hollow-reference-implementation>

<https://github.com/Netflix/hollow>

**Thank
You**